

Dmitry Barsukov

CONTACT INFORMATION	Telegram: @riZZZhik (preferred) GitHub: github.com/riZZZhik LinkedIn: linkedin.com/in/riZZZhik Email: riZZZhik@gmail.com	Residence: Moscow, Russia Open to remote and hybrid full-time jobs
PROFESSIONAL SUMMARY	Machine Learning Engineer with 7+ years of experience, currently focused on high-scale LLM inference optimization. Strong background in production multimodal AI (Text-to-Speech, Speech-to-Text, Computer Vision) and end-to-end deep-learning optimization.	
PROGRAMMING SKILLS	Languages: Python (advanced), Go (advanced), Tritonlang (average), C/C++ (average), CUDA C++ (average) Deep Learning Frameworks: PyTorch, TensorFlow, Keras, scikit-learn Deep Learning Architectures: Transformer, Diffusion, GAN, Speech SOTA, LLM SOTA Optimization frameworks: SGLang, vLLM, TensorRT, TensorRT-LLM, torchcompile, Tritonlang, OpenVINO, ONNX Runtime, LiteRT (f.k.a. TensorFlow Lite) Technical skills: Triton Server, OpenCV, ffmpeg, Torchaudio, WandB, ClearML, Docker, Kubernetes, Git, CI/CD, Observability, Prometheus, Grafana, Linux	
LANGUAGES	Russian (native); English (advanced)	
EMPLOYMENT AND EXPERIENCE	TBank Senior Research Developer, G.E.R.I.	October 2026 - Present
	Optimizing high-scale LLM inference through system design and tuning open-source serving engines.	
	Responsibilities: – System design for distributed LLM serving infrastructure, including KV cache management. – Tuning open-source LLM inference engines (e.g., SGLang, vLLM, TensorRT-LLM) for high-throughput, low-latency serving. – Low-precision quantization (FP8) and other model-level optimizations for in-house LLMs.	
	Achievements: – Applied FP8 quantization with TensorRT-LLM to a custom in-house model, achieving 2x throughput per GPU. – Designed and built a KV cache sharing system, lifting per-GPU throughput by up to 2.2x.	
	Technologies: Python, PyTorch, TensorRT-LLM, FP8 quantization, CUDA, Triton Server, Docker + Kubernetes, Observability	

Victory Group
Machine Learning Team Lead

April 2025 - October 2026

Team lead for a team of 4 data scientists and machine learning engineers, responsible for the development of a multimodal (text, audio, classic) AI systems for real-time text and audio processing.

Responsibilities:

- Leading the team in the research, development and deployment of a multimodal AI systems.
- Mentoring and guiding team members in their professional growth.
- Collaborating with cross-functional teams to ensure successful project delivery.

Achievements:

- Hired and mentored a team of 4 DS and MLE.
- Built from scratch all the processes and infrastructure for the team, including:
 - Model development and deployment processes
 - CI/CD pipelines for model deployment
 - Version control and artifact management
 - Monitoring, logging and alerting
- Under my leadership, the team has successfully developed and deployed:
 - A multimodal AI system combining ASR, TTS, and LLM components for real-time phone conversations that achieve human-level performance.
 - A comprehensive AI solution for automated phone conversation quality assessment and analysis.

MTS AI
Senior Machine Learning Engineer

June 2022 - March 2025

Development of a Text-to-Speech, Speech-to-Text, and ASR services that outperform leading competitors in the Russian language.

Responsibilities:

- Model deployment using Triton Server, Docker, Kubernetes, Python, and Golang.
- Model optimization for performance and resource efficiency:
 - Model architecture changes
 - TensorRT, tritonlang, torchcompile and OpenVINO
 - Model warmup, quantization, sparsity and pruning
- Research, develop, train and fine-tune new model architectures.

Achievements:

- **0.12** p95 latency and **90** RPS for diffusion model on a single **2g.20Gb** A100 instance.
- Established and automated a version-controlled model deployment process using CI/CD, WandB / ClearML, and Artifactory.
- The development process was established following best practices (e.g., CI/CD, code review, documentation, unit testing, changelog, semantic versioning, etc.).
- Created an automated quality and performance testing in a production-like Kubernetes environment.

– Implemented comprehensive observability with monitoring, logging, and alerting.

Technologies:

Python, Golang, PyTorch, tritonlang, CUDA C++, torchcompile, TensorRT, WandB / ClearML, Triton Server, Observability, Docker + Kubernetes, Git + CI/CD

SIRIN

March 2021 - January 2022

Senior Machine Learning Developer

Designed and implemented machine learning service for the automatic opening of car barriers using computer vision.

Achieved **99%** accuracy in recognizing Russian license plates and **90%** accuracy for all other license plates, maintaining a latency of **0.5** on a **4**-core CPU.

Technologies:

Python, PyTorch, OpenCV, Docker + Kubernetes, OpenVINO + Triton Server, Observability (Grafana, Kibana, Prometheus), Git + CI/CD

ITMO University

January 2020 - December 2020

Machine Learning Developer

Designed and implemented service for building facade segmentation, managing everything from data collection and preprocessing to model training, evaluation, and deployment.

Technologies:

Python, TensorFlow + Keras, OpenCV, Docker, Git

SPIIRAS

August 2018 - October 2020

Middle Machine Learning developer

Designed and implemented facial recognition service, managing everything from data collection/generation and preprocessing to model training, evaluation, and deployment.

Technologies:

Python, TensorFlow + Keras, RealSense DepthCamera, OpenCV, Docker, Git

EDUCATION

Higher School of Economics

B.S., Applied Mathematics and Information Science.

GPA: 3.7/4.0

Moscow, Russia (Remote)

September 2023 - Present